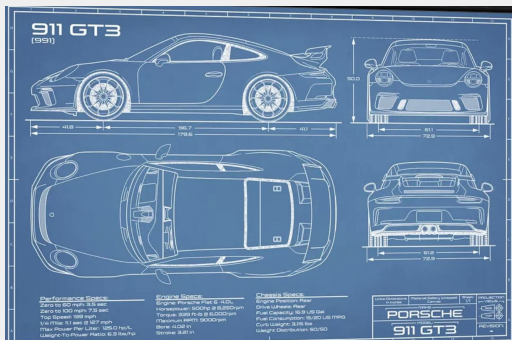


In-repo Open Policy Agent validation of GitOps definitions



Motivation of using “in-repo” OPA versus GateKeeper (blue print check versus conveyor belt check)



- Cheap
- Agnostic of fulfillment
- Validation done locally
- Easy to develop
- Policy validated in PR



- Expensive
- Kubernetes specific
- Requires k8s runtime
- Requires knowledge of CR
- Policy validated at fulfillment phase, post PR merge

Example use case: gitops KCC GCP project IAM definitions (roles)

Branch: master ▼ **cpa-kcc-configs** / [project_vars](#) / [kohlsdev-cpa-inspec](#) / iam-policy-members.yml

 **TKMAM6X** [CPA-2505][add]added opa config and new project vars

1 contributor

19 lines (18 sloc) | 372 Bytes

```
1  ---
2  iamPolicyMembers:
3    usersByEmail:
4      mark.doll@kohls.com:
5        roles:
6        - viewer
7      hamilton.hoover@kohls.com:
8        roles:
9        - editor
10       - bigquery.jobUser
11    groupsByEmail:
12      gcp-et-devops-l3@kohls.com:
13        roles:
14        - viewer
15    serviceAccountsByEmail:
16      gitops-example-2@kohlsdev-cpa-inspec.iam.gserviceaccount.com:
17        roles:
18        - viewer
```

IAM roles are requested for users, groups and service accounts

Can we restrict role allocation automatically before change can be approved ?


Can we define those restriction based on environment of the project (prod vs lle) ?

Can we restrict role allocation based on requestor type: user, group, service account ?

Can we run validation automatically ?

IAM Policy definition (separate policy repository)

Branch: master ▾ cpa-kcc-templates / opa / data / iam-roles.yml

 TKMAM6X added new iam roles defs

1 contributor

51 lines (50 sloc) | 1.5 KB

```
1 ---
2 POLICY_DATA:
3   iam:
4     allowedRoles:
5       # PRODUCTION
6       prd:
7         roles:
8         - editor:
9             allowedForGroups: false
10            allowedForServiceAccounts: true
11            allowedForUsers: false
12        - viewer:
13            allowedForGroups: true
14            allowedForServiceAccounts: true
15            allowedForUsers: true
16        - compute.instanceAdmin.v1:
17            allowedForGroups: true
18            allowedForServiceAccounts: true
19            allowedForUsers: false
20      # LLE
21      lle:
22        roles:
23        - editor:
24            allowedForGroups: true
25            allowedForServiceAccounts: true
26            allowedForUsers: true
27        - viewer:
28            allowedForGroups: true
29            allowedForServiceAccounts: true
```

Allowed IAM roles are whitelisted per environment with additional flags specifying availability per type of requestor

If role is not listed, it is not allowed

This content is managed independently from gitops configuration data by policy maintainers

How policies logic are defined ?

Policy is combination of policy logic (REGO file) and policy data that allows for parameterized evaluation. For example,

Policy data

Branch: master [cpa-kcc-templates](#) / [opa](#) / [data](#) / iam-roles.yml

TKMAM6X added new iam roles defs

1 contributor

51 lines (50 sloc) | 1.5 KB

```
1 ---
2 POLICY_DATA:
3   iam:
4     allowedRoles:
5       # PRODUCTION
6       prd:
7         roles:
8           - editor:
9             allowedForGroups: false
10             allowedForServiceAccounts: true
11             allowedForUsers: false
12           - viewer:
13             allowedForGroups: true
14             allowedForServiceAccounts: true
15             allowedForUsers: true
16           - compute.instanceAdmin.v1:
17             allowedForGroups: true
18             allowedForServiceAccounts: true
19             allowedForUsers: false
20   # LLE
21   lle:
22     roles:
23       - editor:
24         allowedForGroups: true
25         allowedForServiceAccounts: true
26         allowedForUsers: true
27       - viewer:
28         allowedForGroups: true
29         allowedForServiceAccounts: true
```

REGO file

Branch: master [cpa-kcc-templates](#) / [opa](#) / [rego](#) / iam-roles.rego

Find file Copy path

TKMAM6X [CPA-2505] added opa definitions

938898b 8 hours ago

1 contributor

54 lines (44 sloc) | 2.06 KB

Raw Blame History

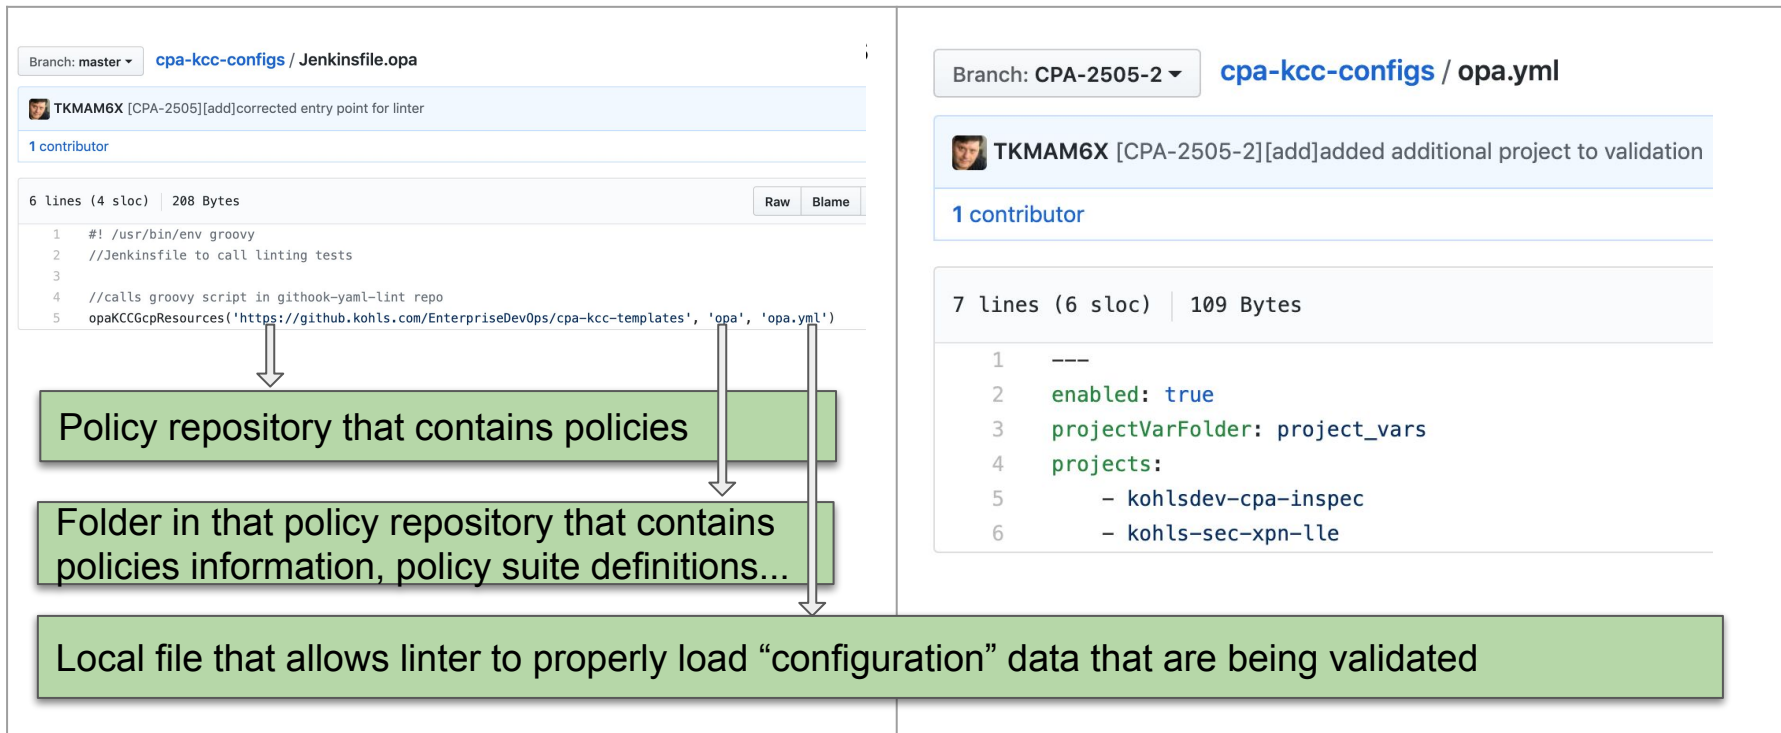
```
1 package kohlts.gitops.kcc.gcp.project.iam.roles
2 default allow = false
3
4 environment := input.project.labels["environment-type"]
5
6 # users with associated IAM roles
7 users := input.iamPolicyMembers.usersByEmail
8 user_roles := [item | item := users[_].roles]
9 # flatten it
10 requested_user_roles := [item | item := user_roles[_][_]]
11 # all allowed user IAM roles in given environment
12 allowed_user_roles := [role | item := data.POLICY_DATA.iam.allowedRoles[environment].roles[index][role]; data.POLICY_DATA.iam.allowedRoles[environment].roles[index][role]]
13
14 # groups with associated IAM roles
15 groups := input.iamPolicyMembers.groupsByEmail
16 group_roles := [item | item := groups[_].roles]
17 # flatten it
18 requested_group_roles := [item | item := group_roles[_][_]]
19 # all allowed group IAM roles in given environment
20 allowed_group_roles := [role | item := data.POLICY_DATA.iam.allowedRoles[environment].roles[index][role]; data.POLICY_DATA.iam.allowedRoles[environment].roles[index][role]]
21
22 # Service Accounts with associated IAM roles
23 sas := input.iamPolicyMembers.serviceAccountsByEmail
24 sa_roles := [item | item := sas[_].roles]
25 # flatten it
26 requested_sa_roles := [item | item := sa_roles[_][_]]
27 # all allowed Service Account IAM roles in given environment
28 allowed_sa_roles := [role | item := data.POLICY_DATA.iam.allowedRoles[environment].roles[index][role]; data.POLICY_DATA.iam.allowedRoles[environment].roles[index][role]]
29
30 # Allow if "all and each" of requested "kind" roles are present in allowed_kinds_roles
31 allow {
32   all_in(requested_user_roles, allowed_user_roles)
33   all_in(requested_group_roles, allowed_group_roles)
34   all_in(requested_sa_roles, allowed_sa_roles)
35 }
36
37 # ----- Library routines to implement all_in() and in_list() validation -----
38 all_in(items, list) {
39   not_any_not_in(items, list)
40 }
41
42 not_any_not_in(items, list) {
43   l = [item | item := not_in_list(items[index], list)]
44   count(l) == 0
45 }
46
47 not_in_list(item, list) {
48   not_in_list(item, list)
49 }
50
51 in_list(item, list) {
52   list[_] == item
53 }
```

How OPA policy is evaluated

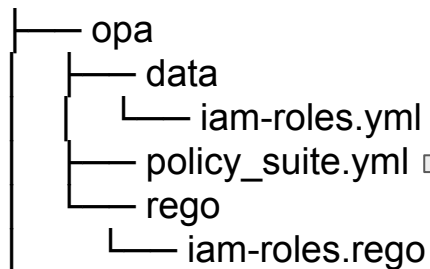
- Policy is evaluated by combining policy logic (REGO) with policy data and “evaluating” it against context of gitops data definitions known as “input”
- Policy evaluation results in JSON output that will be interpreted as “allowed” or “not allowed”
- Policy evaluator is implemented as GitHub linter associated with repository that needs to be checked

How OPA policies validation is enabled ?

Policy is enabled by configuring **Jenkinsfile.opa** and **opa.yml** in git ops **configuration repository** to “link” gitops definitions with policy definitions located in **different repository**. Jenkinsfile.opa file triggers opa linter and informs it about location of the policies and also provides information on how to prepare input data.



Sample content of policy repo referred to from config repo (policy_suite.yml is expected to be present)



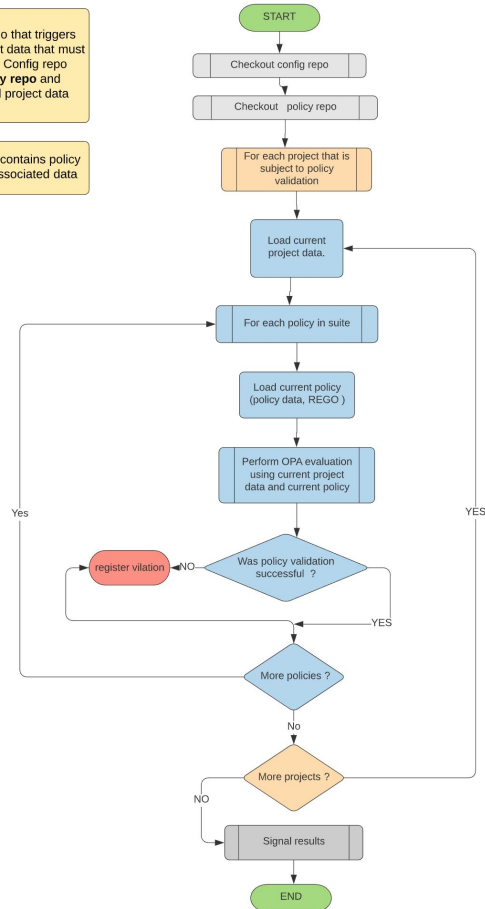
policy_suite.yml

```
1  ---
2  policies:
3    - name: Users Groups and Service Accounts are only given allowed roles
4      active: true
5      policyDataFile: data/iam-roles.yml
6      policyRegoFile: rego/iam-roles.rego
7      package: "kohls.gitops.kcc.gcp_project.iam.roles"
8
9
```



Policy validator logical workflow


Config repo is the github repo that triggers validation and it contains project data that must be checked against policies. Config repo contains reference to **policy repo** and information on how load project data

Policy repo is the github repo contains policy suite that defines policy and associated data





Example policy validation run. GitOps view





**All checks have passed**

6 successful checks



[Hide all checks](#)

 **Open Policy Agent Validation** — OPA Validation completed



[Details](#)

 **PyLinter** — Pylint passed



[Details](#)

 **Schema Validation** — Schema validation passed


[Details](#)

 **Shell Linter** — SHELL linting passed

[Details](#)

 **YAML Linter** — YAML linting passed

[Details](#)

**This branch has no conflicts with the base branch**

Merging can be performed automatically.

Merge pull request

▼

You can also [open this in GitHub Desktop](#) or view [command line instructions](#).

Example policy validation run. Details (Jenkins view)

→ ↻ ↗ iaas2-dev1-jenkins.kohls.com:8443/blue/organizations/jenkins/githubhooks%2FOPA_Validation%2Fcpa-kcc-configs/detail/CPA-2505-2/3/pipeline

Apps Kohl's Links KOHLS APPS myKohls GCP console Good Easy Reading iaas2->paas Training Docs Jenkins GitHubRepos

- ✓ > opa_validation/opa_linter.py — Load a resource file from a shared library
- ✓ > opa_validation/requirements.txt — Load a resource file from a shared library
- ✓ > opa_validation/opa_run.py — Load a resource file from a shared library
- ✓ > requirements.txt — Write file to workspace
- ✓ > main_script.py — Write file to workspace
- ✓ > runner_script.py — Write file to workspace
- ✓ > pip3 install -r requirements.txt --user --proxy=<http://proxy-gcp-central.kohls.com:8080> — Shell Script
- ✓ > Shell Script
- ✓ ▼ cat /tmp/opa.log — Shell Script

```
1 + cat /tmp/opa.log
2 2020-07-14 16:22:28,855 INFO: Loaded OPA validation config
3 2020-07-14 16:22:28,856 INFO: Processing project data for kohls-sec-xpn-lle
4 2020-07-14 16:22:29,042 INFO: Processing policy 'Users Groups and Service Accounts are only given allowed roles'
5 2020-07-14 16:22:31,476 INFO: OPA evaluation completed
6 2020-07-14 16:22:31,476 INFO: Result set. START
7 2020-07-14 16:22:31,476 INFO: {
8   "result": [
9     {
10      "expressions": [
11        {
12          "value": {
13            "allow": true,
14            "allowed_group_roles": [
15              "editor",
16              "viewer"
17            ],
18            "allowed_sa_roles": [
19              "editor",
20              "viewer",
21              "bigquery.jobUser",
22              "compute.networkUser",
23              "compute.networkUser",
24              "container.hostServiceAgentUser",
25              "compute.networkUser"
26            ],
27            "allowed_user_roles": [
28              "editor",
29              "viewer",
30              "bigquery.jobUser"
31            ],
32            "environment": "lle",
```

Interpretation of OPA policy validation response

```
2020-07-14 16:22:28,855 INFO: Loaded OPA validation config
2020-07-14 16:22:28,856 INFO: Processing project data for kohls-sec-xpn-lle
2020-07-14 16:22:29,042 INFO: Processing policy 'Users Groups and Service Accounts are only given allowed roles'
2020-07-14 16:22:31,476 INFO: OPA evaluation completed
2020-07-14 16:22:31,476 INFO: Result set. START
2020-07-14 16:22:31,476 INFO: Result: {
  "result": {
    {
      "expressions": [
        {
          "value": {
            "allow": true,
            "allowed_group_roles": [
              "editor",
              "viewer"
            ],
            "allowed_sa_roles": [
              "editor",
              "viewer",
              "bigquery.jobUser",
              "compute.networkUser",
              "compute.networkUser",
              "container.hostServiceAgentUser",
              "compute.networkUser"
            ],
            "allowed_user_roles": [
              "editor",
              "viewer",
              "bigquery.jobUser"
            ]
          },
          "environment": "lle",
          "group_roles": [],
          "requested_group_roles": [],
          "requested_sa_roles": [
            "container.hostServiceAgentUser",
            "compute.networkUser",
            "compute.networkUser",
            "compute.networkUser"
          ],
          "requested_user_roles": [],
          "sa_roles": [
            "container.hostServiceAgentUser",
            "compute.networkUser"
          ],
          "user_roles": [
            "compute.networkUser"
          ],
          "sas": {
            "938246310506-compute@developer.gserviceaccount.com": {
              "roles": [
                "compute.networkUser"
              ]
            },
            "938246310506@cloudservices.gserviceaccount.com": {
              "roles": [
                "compute.networkUser"
              ]
            },
            "service-938246310506@container-engine-robot.iam.gserviceaccount.com": {
              "roles": [
                "container.hostServiceAgentUser",
                "compute.networkUser"
              ]
            }
          },
          "user_roles": []
        },
        {
          "text": "data.kohls.gitops.kcc.gcp_project.iam.roles",
          "location": {
            "row": 1,
            "col": 1
          }
        }
      ]
    }
  }
}
```

2020-07-14 16:22:31,476 INFO: Result set. END
2020-07-14 16:22:31,476 INFO: >>>>>>> 'Users Groups and Service Accounts are only given allowed roles' ALLOWED
2020-07-14 16:22:31,482 INFO: Processing project data for kohls-sec-xpn-lle

Indicates what data are used for input and what policy is used

"allow: true"
Means policy is allowed.

Debug info: shows roles (user, group, SA) allowed by policy definition

Debug info: shows roles (user, group, SA) requested by gitops config

References

Config repo (data that are being validated)

<https://github.kohls.com/EnterpriseDevOps/cpa-kcc-configs/tree/CPA-2505-2>

Repo that contains policy definitions

<https://github.kohls.com/EnterpriseDevOps/cpa-kcc-templates>

Validator (linter) code

https://github.kohls.com/EnterpriseDevOps/linter_dev

Workflow diagram (policy validation process)

<https://app.lucidchart.com/documents/view/48648ba7-481c-4833-9fb4-bb1be181e03c>

<https://www.openpolicyagent.org/>