


Services and Networking

What is Networking?



Simply put, Networking is when one computer can talk to another. This communication happens in many diverse ways but you can think of it like a FedEx delivery.

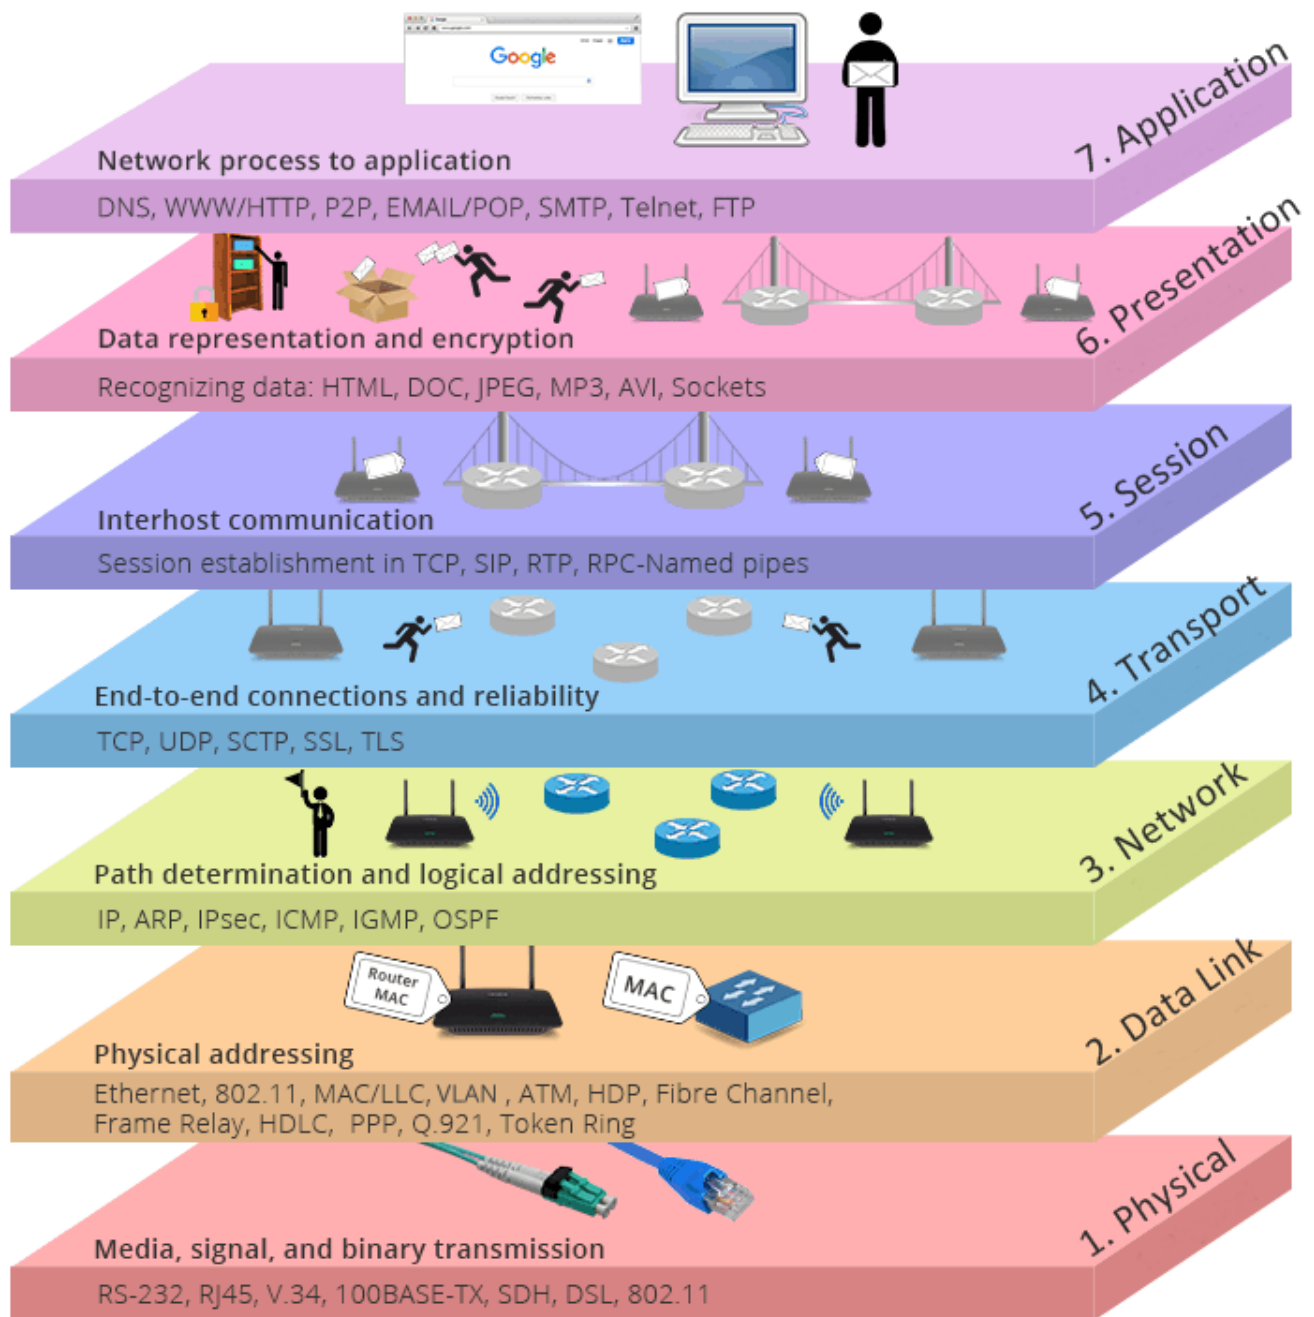
Essentially all devices connected to the internet have an IP address, which is like your house address.

They also have their own MAC address, which is the like name of the home owner.

Sending a package to another person is considered creating "traffic" over the internet. Packages or **packets** in networking, traverse the internet to send data to your chosen destination. The packet will have your IP address and MAC address, the same way a FedEx package would have your name and home address.


Networking includes all of the protocols or rules used on the internet to communicate from point a to point b. Some of those protocols regulate services like DNS, HTTP, or HTTPS traffic.

The OSI Model



Understanding Connectivity Between Pods

A pod consists of one or more containers that are collocated on the same host, and are configured to share a network stack and other resources such as volumes. All pods on a container can reach one another on a localhost. If I have a container running nginx and listening on port 80 and another container running prometheus, the second container can connect to the first via <http://localhost>.



Here we have two separate pods running container 1 and container 2. The eth0 port is the physical interface that then uses docker0 as a bridge for the containers.

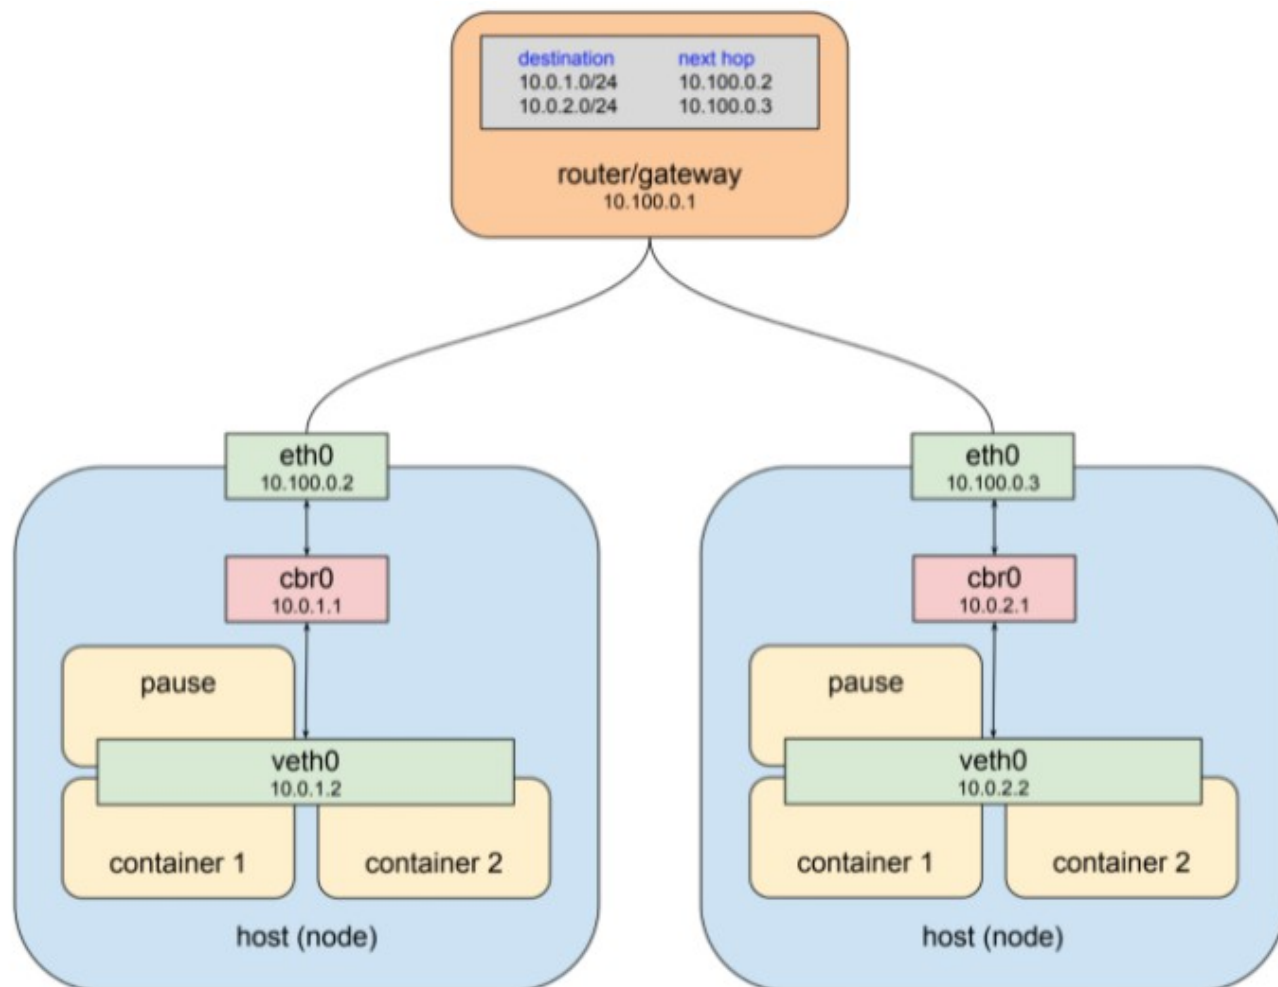
Notice that both virtual ethernet interfaces have separate IP's but can communicate by default because they're in the same subnet.

Kubernetes implements this pattern by creating a special container for each pod whose only purpose is to provide a network interface for the other containers. If you ssh in to a node that has pods scheduled on it and run `docker ps` you will see at least one container that was started with the `pause` command. The `pause` command suspends the current process until a signal is received so these containers do nothing at all except sleep until Kubernetes sends them a SIGTERM.

```
docker@minikube:~$ docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS
27adb36fac47   8d147537fb7d                       "/coredns -conf /etc..." Less than a second ago Up
527af1f04238   6120bd723dce                       "/usr/local/bin/kube..." 1 second ago   Up
f99e4db4b3ba   6e38f40d628d                       "/storage-provisioner"     1 second ago   Up
3c6e289dcd47   k8s.gcr.io/pause:3.5               "/pause"                 1 second ago   Up
cd0895f4d2ee   k8s.gcr.io/pause:3.5               "/pause"                 1 second ago   Up
765b3e776427   k8s.gcr.io/pause:3.5               "/pause"                 1 second ago   Up
```

Despite this lack of activity the "pause" container is the heart of the pod, providing the virtual network interface that all the other containers will use to communicate with each other and the outside world.

The private networks 10.100.0.2 and 10.100.0.3 are for the two instances in this example and our router is 10.100.0.1. Given this setup, each instance can communicate with the other on eth0. The pods/clusters are hanging off a bridge on a different network entirely, one that is virtual and exists only on a specific node.



The easiest way to think about ethernet is the one on your laptop. These eth0 ports exist physically and also in virtualized software to provide a connection from the device you're using to another server, cloud router, or physical router responsible for directing internet traffic.



A Brief Introduction to Subnetting

Subnets				CIDR/IPv4 Cheat Sheet		
CIDR	Subnet Mask	# of Addresses	Wildcard	Classful Ranges		
/0	0.0.0.0	4,294,967,296	255.255.255.255	A 0.0.0.0 - 127.255.255.255		
/1	128.0.0.0	2,147,483,648	127.255.255.255	B 128.0.0.0 - 191.255.255.255		
/2	192.0.0.0	1,073,741,824	63.255.255.255	C 192.0.0.0 - 223.255.255.255		
/3	224.0.0.0	536,870,912	31.255.255.255	D 224.0.0.0 - 239.255.255.255		
/4	240.0.0.0	268,435,456	15.255.255.255	E 240.0.0.0 - 255.255.255.255		
/5	248.0.0.0	134,217,728	7.255.255.255	Reserved Ranges		
/6	252.0.0.0	67,108,864	3.255.255.255	RFC 1918 10.0.0.0 - 10.255.255.255		
/7	254.0.0.0	33,554,432	1.255.255.255	Localhost 127.0.0.0 - 127.255.255.255		
/8	255.0.0.0	16,777,216	0.255.255.255	RFC 1918 172.16.0.0 - 172.31.255.255		
/9	255.128.0.0	8,388,608	0.127.255.255	RFC 1918 192.168.0.0 - 192.168.255.255		
/10	255.192.0.0	4,194,304	0.63.255.255	CIDR notation		
/11	255.224.0.0	2,097,152	0.31.255.255	Classless interdomain routing (CIDR) notation is a compact representation of an IP address and its' associated routing prefix. It's expressed as a / followed by a number (e.g. /0 or /10).		
/12	255.240.0.0	1,048,576	0.15.255.255			
/13	255.248.0.0	524,288	0.7.255.255			
/14	255.252.0.0	262,144	0.3.255.255			
/15	255.254.0.0	131,072	0.1.255.255			
/16	255.255.0.0	65,536	0.0.255.255	VLSM		
/17	255.255.128.0	32,768	0.0.127.255	CIDR is based on the variable-length subnet masking (VLSM) technique, which allows the specification of arbitrary-length prefixes.		
/18	255.255.192.0	16,384	0.0.63.255			
/19	255.255.224.0	8,192	0.0.31.255			
/20	255.255.240.0	4,096	0.0.15.255			
/21	255.255.248.0	2,048	0.0.7.255			
/22	255.255.252.0	1,024	0.0.3.255	Decimal to Binary		
/23	255.255.254.0	512	0.0.1.255	Subnet Mask		Wildcard
/24	255.255.255.0	256	0.0.0.255	0	0000 0000	255 1111 1111
/25	255.255.255.128	128	0.0.0.127	128	1000 0000	127 0111 1111
/26	255.255.255.192	64	0.0.0.63	192	1100 0000	63 0011 1111
/27	255.255.255.224	32	0.0.0.31	224	1110 0000	31 0001 1111
/28	255.255.255.240	16	0.0.0.15	240	1111 0000	15 0000 1111
/29	255.255.255.248	8	0.0.0.7	248	1111 1000	7 0000 0111
/30	255.255.255.252	4	0.0.0.3	252	1111 1100	3 0000 0011
/31	255.255.255.254	2	0.0.0.1	254	1111 1110	1 0000 0001
/32	255.255.255.255	1	0.0.0.0	255	1111 1111	0 0000 0000

Hands On - Connect a Pod to a Service

I. Setup your deployment and Verify

The first step requires us deploy the following manifest. This will setup a pod and service that allows network communication on port 80.

Service

```
kind: Service
apiVersion: v1
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Pod

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Create your service and pod, then apply them to the cluster. Once done, verify that both are running.

```
dominickhrndz314@cloudshell:~$ nano yourname-service.yaml
dominickhrndz314@cloudshell:~$ nano yourname-pod.yaml
dominickhrndz314@cloudshell:~$ kubectl apply -f yourname-service.yaml
service/nginx-service created
dominickhrndz314@cloudshell:~$ kubectl apply -f yourname-pod.yaml
deployment.apps/nginx-deployment created
dominickhrndz314@cloudshell:~$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-deployment-5d59d67564-dwsc6	0/1	ContainerCreating	0	7s
pod/nginx-deployment-5d59d67564-jpf8z	1/1	Running	0	7s
pod/nginx-deployment-5d59d67564-vbngt	0/1	ContainerCreating	0	7s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	44m
service/nginx-service	ClusterIP	10.111.255.3	<none>	80/TCP	15s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/nginx-deployment	1/3	3	1	7s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/nginx-deployment-5d59d67564	3	3	1	7s

```
dominickhrndz314@cloudshell:~$
```


II. Identify the pod IP's and verify connectivity between pods

Lets first identify the IP's associated with our deployments. To this we run a new command: `kubectl get po -l app=nginx -o wide`. Then we can check connectivity of our pods by running the `ping` command.

```
dominickhrndz314@cloudshell:~$ kubectl get po -l app=nginx -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP
nginx-deployment-5d59d67564-dwsc6   1/1     Running   0           3m44s 172.17.0.4
nginx-deployment-5d59d67564-jpf8z   1/1     Running   0           3m44s 172.17.0.3
nginx-deployment-5d59d67564-vbngt   1/1     Running   0           3m44s 172.17.0.5
dominickhrndz314@cloudshell:~$
```

Now lets SSH into our pod and test communication between pods. We can issue the command `kubectl exec -it <name of pod> - bin/bash` to login to the pod of our choice. First run a ping to a destination pod.

```
dominickhrndz314@cloudshell:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-5d59d67564-dwsc6   1/1     Running   0           9m41s
nginx-deployment-5d59d67564-jpf8z   1/1     Running   0           9m41s
nginx-deployment-5d59d67564-vbngt   1/1     Running   0           9m41s
dominickhrndz314@cloudshell:~$ kubectl exec -it nginx-deployment-5d59d67564-dwsc6 -- /bin/bash
root@nginx-deployment-5d59d67564-dwsc6:/# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
7: eth0@if9: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:ac:11:00:04 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.4/16 brd 172.17.255.255 scope global eth0
        valid_lft forever preferred_lft forever
root@nginx-deployment-5d59d67564-dwsc6:/# ping 172.17.0.3
PING 172.17.0.3 (172.17.0.3): 48 data bytes
56 bytes from 172.17.0.3: icmp_seq=0 ttl=64 time=0.117 ms
56 bytes from 172.17.0.3: icmp_seq=1 ttl=64 time=0.077 ms
^C--- 172.17.0.3 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max/stddev = 0.077/0.097/0.117/0.000 ms
root@nginx-deployment-5d59d67564-dwsc6:/# ping 172.17.0.5
PING 172.17.0.5 (172.17.0.5): 48 data bytes
56 bytes from 172.17.0.5: icmp_seq=0 ttl=64 time=0.131 ms
56 bytes from 172.17.0.5: icmp_seq=1 ttl=64 time=0.081 ms
^C--- 172.17.0.5 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max/stddev = 0.081/0.106/0.131/0.025 ms
root@nginx-deployment-5d59d67564-dwsc6:/#
```

If your ping is successful you have completed this lesson for pod to pod communications.

